

Building infrastructure orchestration software

Patryk Orwat

14.01.2025



How to manage 1M+ resources in a centralized fashion?

Order of battle

- Infra management goal
- K8S concepts
- Hykube playground
- Crossplane case
- Other alternatives
- Lessons learnt

Intro

Why K8S?

- API interface for resource modelling is well defined (better than HATEOAS)
- Developers, operations, testers are familiarized with `kubectl`
- There's wide ecosystem around K8S (Helm, Crossplane, Velero, K3S, Argo, ...)

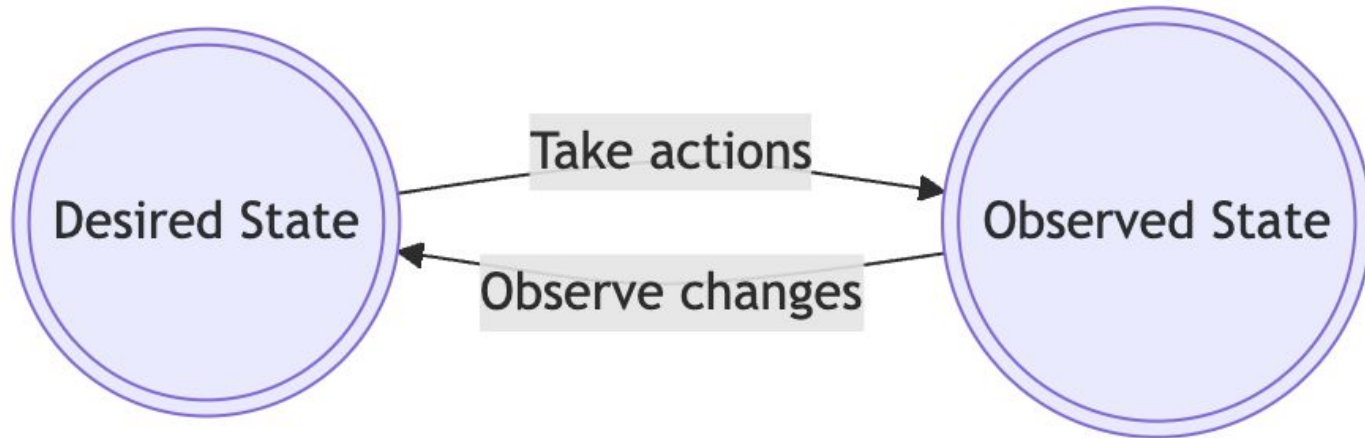
K8S for infra pain points - scalability

- As stated in the documentation for version 1.30.0, it can support
 - up to: 5K nodes
 - 150K pods
 - 300K total containers
- Even though higher numbers are possible, e.g., done by Google, it's still an order of magnitude lower than hyperscale level clusters with 100,000+ nodes.

K8S for infra pain points - containerness bound

- K8S was created to manage containers
- Container-related resources are in the core K8S code
- Reconciliation loop...

K8S for infra pain points - reconciliation loop



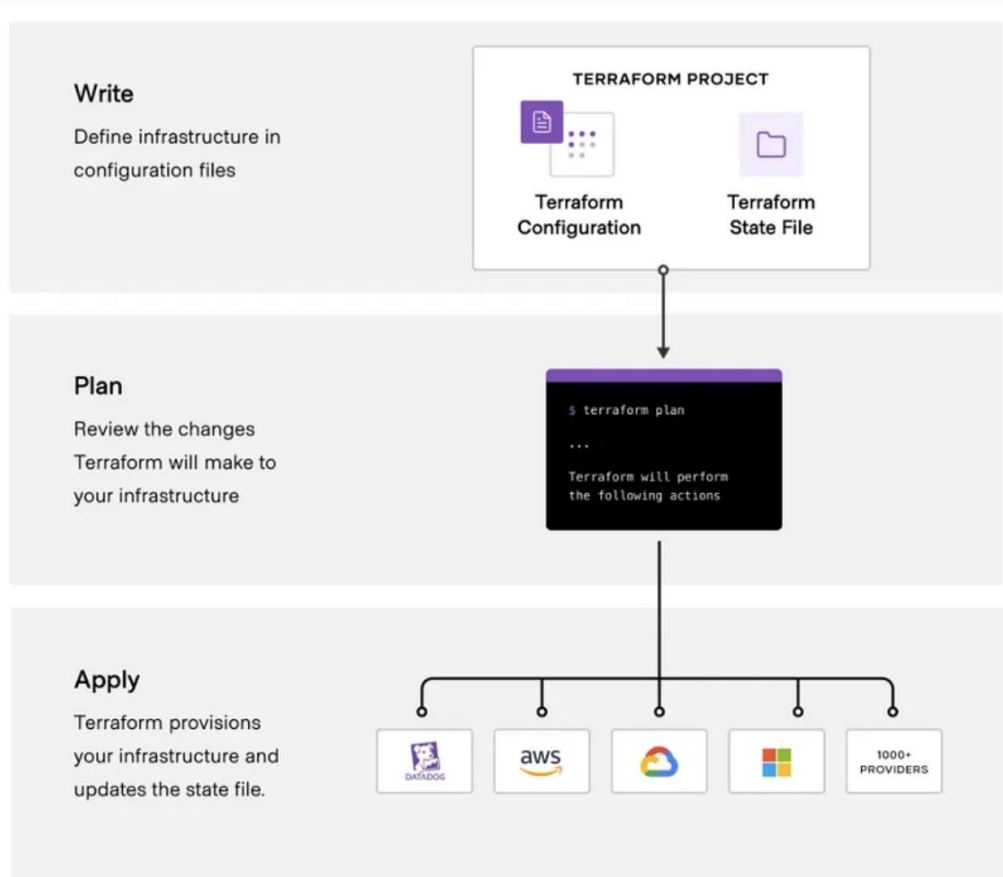
K8S for infra pain points - reconciliation loop

- Loops allow for intent-based management
- Impossible to aggregate reconciliation requests
- Watching mechanism takes up resources

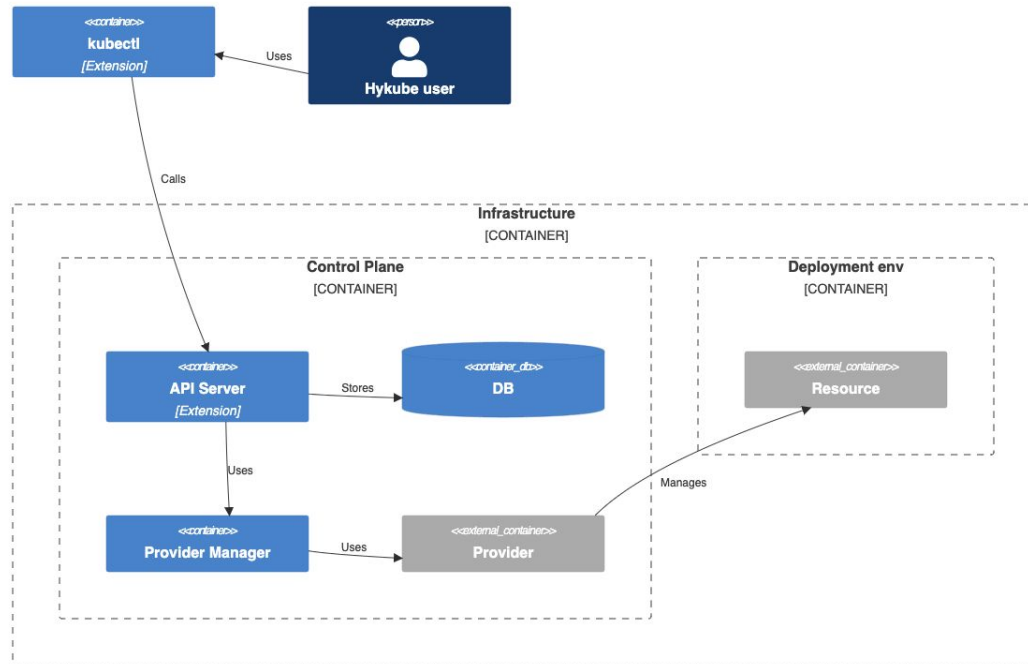
Attempt 1 - hykubc

Hykube main points

- Let's walk around reconciliation loop
- Add plan (generate an execution plan) and apply (execute plan actions) commands, see Terraform
- Wrap it with K8S API



Hykube Architecture



K8S - kubectl plugins

Adding new commands to plan and apply is easy

Example plugin

```
#!/bin/bash

# optional argument handling
if [[ "$1" == "version" ]]
then
    echo "1.0.0"
    exit 0
fi

# optional argument handling
if [[ "$1" == "config" ]]
then
    echo "$KUBECONFIG"
    exit 0
fi

echo "I am a plugin named kubectl-foo"
```

->

To use a plugin, make the plugin executable:

```
sudo chmod +x ./kubectl-foo
```

and place it anywhere in your `PATH` :

```
sudo mv ./kubectl-foo /usr/local/bin
```

->

You may now invoke your plugin as a `kubectl` command:

```
kubectl foo
```

```
I am a plugin named kubectl-foo
```

All args and flags are passed as-is to the executable:

```
kubectl foo version
```

```
1.0.0
```

K8S - API server vs Operator

Custom API server	Operator pattern (CRDs)
API server is used (almost) directly by kubectl, via apiregistration.k8s.io/v1/APIService	Default K8S API server is used
Flexibility on resource storage	Making use of apiextension API server
Practically, only Go can be used, with forking sample-apiserver	Many languages possible, with Operator Framework/kubebuilder/kopf
You can implement own reconciliation loop	Making use of webhooks and Controllers

Defining resources

K8S resource - definition*

```
// +genclient
// +k8s:deepcopy-gen:interfaces=k8s.io/apimachinery/pkg/runtime.Object
// +k8s:prerelease-lifecycle-gen:introduced=1.0
// +k8s:prerelease-lifecycle-gen:removed=1.10

type Provider struct {
    metav1.TypeMeta   `json:",inline"`
    metav1.ObjectMeta `json:"metadata,omitempty" protobuf:"bytes,1,opt,name=metadata"`

    Spec      ProviderSpec `json:"spec,omitempty" protobuf:"bytes,2,opt,name=spec"`
    Status    string       `json:"status,omitempty" protobuf:"bytes,3,opt,name=status"`
    Filename  string       `json:"filename,omitempty" protobuf:"bytes,4,opt,name=filename"`
}
```

* This is only Go logic; CRDs still require yaml definition (see [YAML](#) and [Go struct](#) example in Crossplane)

K8S resource - Defining a scheme

```
var (  
    // SchemeBuilder is the scheme builder with scheme init functions to run for this API package  
    SchemeBuilder = runtime.NewSchemeBuilder(addKnownTypes)  
    // AddToScheme is a common registration function for mapping packaged scoped group & version keys to a scheme  
    AddToScheme = SchemeBuilder.AddToScheme  
)  
  
// Adds the list of known types to the given scheme.  
✓ func addKnownTypes(scheme *runtime.Scheme) error {  
    scheme.AddKnownTypes(SchemeGroupVersion,  
        &Provider{},  
        &ProviderList{},  
  
        &Plan{},  
        &PlanList{},  
    )  
    return nil  
}
```

K8S resource - Using a scheme in API server

```
// New returns a new instance of HykubeServer from the given config.
func (c completedConfig) New() (*HykubeServer, error) {
    genericServer, err := c.GenericConfig.New("hykube", genericapiserver.NewEmptyDelegate())
    if err != nil {
        return nil, err
    }

    s := &HykubeServer{
        GenericAPIServer: genericServer,
    }

    apiGroupInfo := genericapiserver.NewDefaultAPIGroupInfo(hykube.GroupName, Scheme, metav1.ParameterCodec, Codecs)

    v1alpha1storage := map[string]rest.Storage{}
    v1alpha1storage["providers"] = hykuberegistry.RESTInPeace(providerstorage.NewREST(Scheme, c.GenericConfig.RESTOptionsGetter))
    v1alpha1storage["plans"] = hykuberegistry.RESTInPeace(planstorage.NewREST(Scheme, c.GenericConfig.RESTOptionsGetter))
    apiGroupInfo.VersionedResourcesStorageMap["v1alpha1"] = v1alpha1storage

    if err := s.GenericAPIServer.InstallAPIGroup(&apiGroupInfo); err != nil {
        return nil, err
    }

    return s, nil
}
```

K8S resource - Using a scheme in API client

```
crossplaneSchema := runtime.NewScheme()
_ = crossapiextensionsv1.AddToScheme(s)

crossplaneClient, err := client.New(cfg, client.Options{Scheme: crossplaneSchema})
if err != nil {
    return errors.Wrap(err, "cannot create client")
}
```

```
allRevisions := &crossapiextensionsv1.CompositionRevisionList{}
err = c.crossplaneClient.List(ctx, allRevisions, client.InNamespace(ns.GetName()))
```

What I managed to do?

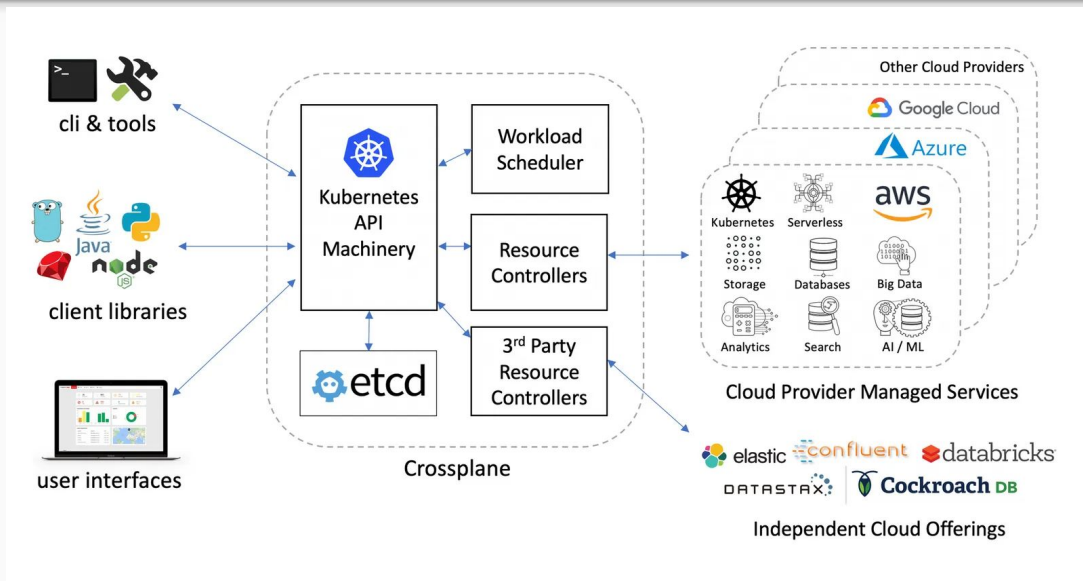
Implemented fetching of a
Terraform plugins +
creating CRDs

```
(base) patrykorwat@Patryks-MacBook-Pro ~ % kubectl get crd | head -n 10
NAME                                     CREATED AT
aws-accessanalyzer-analyzers.aws.hykube.io 2024-09-12T15:06:22Z
aws-accessanalyzer-archive-rules.aws.hykube.io 2024-09-12T15:04:36Z
aws-account-alternate-contacts.aws.hykube.io 2024-09-12T15:05:28Z
aws-account-primary-contacts.aws.hykube.io 2024-09-12T15:07:26Z
aws-account-regions.aws.hykube.io 2024-09-12T15:04:31Z
aws-acm-certificate-validations.aws.hykube.io 2024-09-12T15:05:22Z
aws-acm-certificates.aws.hykube.io 2024-09-12T15:05:34Z
aws-acmpca-certificate-authority-certificates.aws.hykube.io 2024-09-12T15:04:08Z
aws-acmpca-certificate-authoritys.aws.hykube.io 2024-09-12T15:06:40Z
(base) patrykorwat@Patryks-MacBook-Pro ~ % kubectl get crd | wc -l
1372
(base) patrykorwat@Patryks-MacBook-Pro ~ %
```

```
(base) patrykorwat@Patryks-MacBook-Pro ~ % kubectl describe crd aws-s3-bucket | head -n 40
Name:      aws-s3-bucket-accelerate-configurations.aws.hykube.io
Namespace:
Labels:    <none>
Annotations: <none>
API Version: apixtensions.k8s.io/v1
Kind:       CustomResourceDefinition
Metadata:
  Creation Timestamp: 2024-09-12T15:03:53Z
  Generation:        1
  Resource Version:   52675
  UID:                4583d2e2-bdc7-4bf5-b511-e014c96e55d3
Spec:
  Conversion:
    Strategy: None
  Group:      aws.hykube.io
  Names:
    Kind:      aws-s3-bucket-accelerate-configuration
    List Kind:  aws-s3-bucket-accelerate-configurationList
    Plural:     aws-s3-bucket-accelerate-configurations
    Singular:   aws-s3-bucket-accelerate-configuration
  Scope:       Namespaced
  Versions:
    Name: v1
  Schema:
    openAPIV3Schema:
      Properties:
        Bucket:
          Type: string
          X - Kubernetes - Preserve - Unknown - Fields: true
        expected_bucket_owner:
          Type: string
          X - Kubernetes - Preserve - Unknown - Fields: true
        Id:
          Type: string
          X - Kubernetes - Preserve - Unknown - Fields: true
        Status:
          Type: string
          X - Kubernetes - Preserve - Unknown - Fields: true
      Required:
        bucket
```

Attempt 2 - Crossplane

Crossplane - brief overview

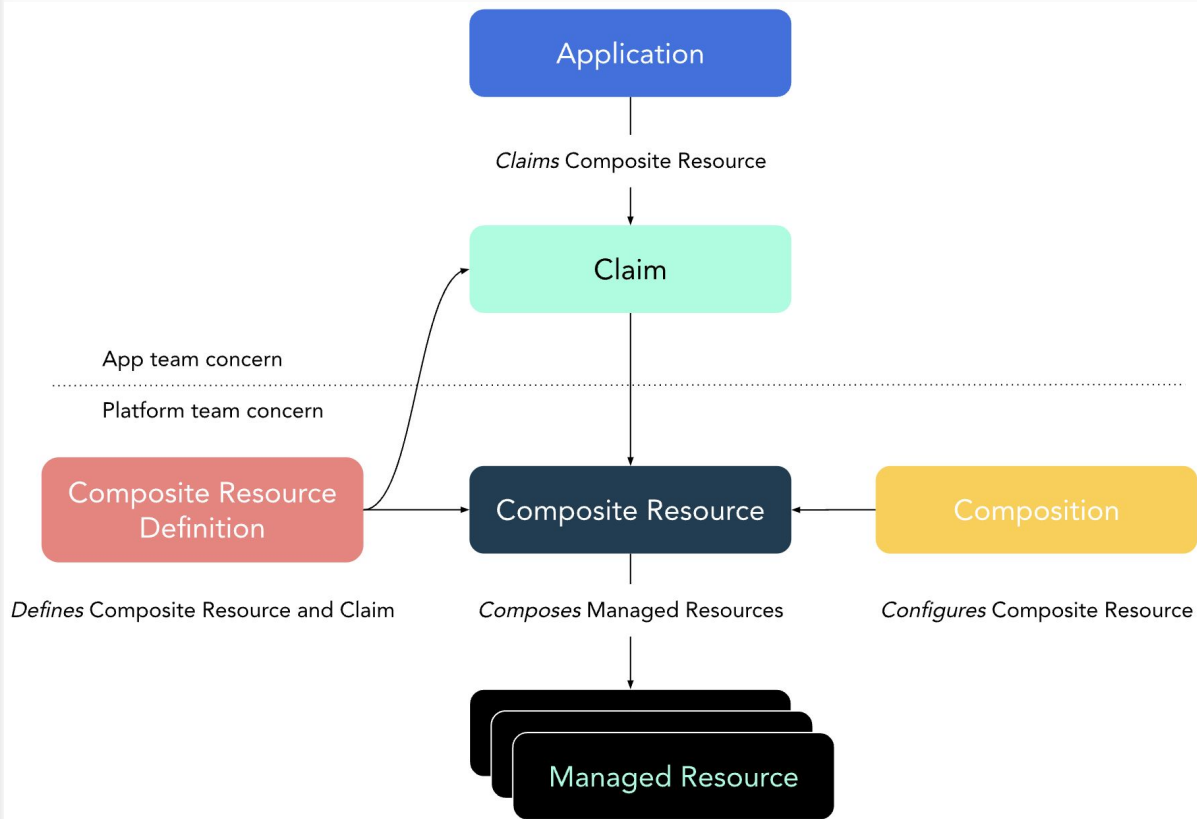


Ref: <https://blog.upbound.io/introducing-crossplane-open-source-multicloud-control-plane>

Crossplane - Composite Resource Definitions

Although

<https://github.com/crossplane/crossplane/discussions/4178>



Ref: <https://docs.crossplane.io/latest/concepts/claims/>

Crossplane - summary

- Tool to manage external resources with K8S interface
- Adds composite resources
- Has okay-ish support of cloud resources (thanks to Terraform)

Crossplane - “reconciliation” loop

For polled resources (in core AND provider):

- *--poll-interval* - check managed resources for spec changes (default: 1 min)
- *--sync-interval* - check external resources and sync with their state (default: 1 hour)

For watched resources:

- *--enable-realtime-compositions* - turn on XRD watching (turn off polling)
- *--enable-watches* - per provider, turn on watching for managed resources

For all resources:

- *--max-reconcile-rate* - rate limiter for REST client & Controller config for MaxConcurrentReconciles; to be configured for core AND per provider

Terraform/Crossplane vs scale

- Core+providers don't support batch operations
 - Reconciliation loop polls for each resource changes, not in batches
- You must model a batch resource - custom resources
 - Reconciliation loop will pool a batched resource, that will perform optimal-ish requests
 - Max default K8S resource size is 1 MB (ETCD limit)

Crossplane open issues

- No clear path on managing 1K+ of resources, little talk about it
- Creating XRDs to model batching is painful
 - Kro makes XRDs creation simpler but is not a solution <https://github.com/awslabs/kro>
 - Need a dedicated model - doesn't seem hard to do

Lessons learnt

Intent-driven orchestration vs workflows

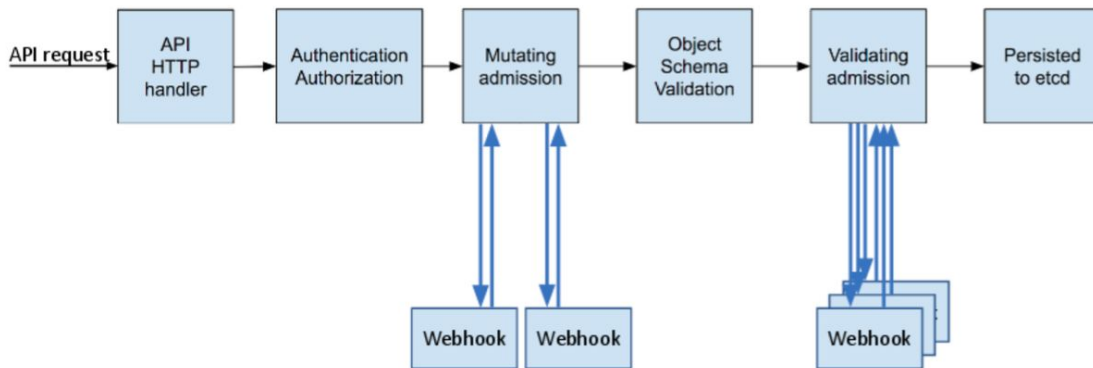
- Intent-driven orchestration:
 - Works well for small quantity of objects (no of batches?)
- Workflows:
 - Define a flow to orchestrate resources

As of today: is there any better Golang OSS than Crossplane + Argo Workflow?

2025 plan - create a resource wrapper + sample workflow

Custom API server vs Operator

- Custom API server works if you:
 - Want to persist data in a separate place
 - Have no reason to make use of reconciliation loop and don't wanna use mutating/validating webhooks



Ref <https://kubernetes.io/blog/2019/03/21/a-guide-to-kubernetes-admission-controllers/>

External calls always slow things down...

Golang with infra future

Infrastructure from/with Code

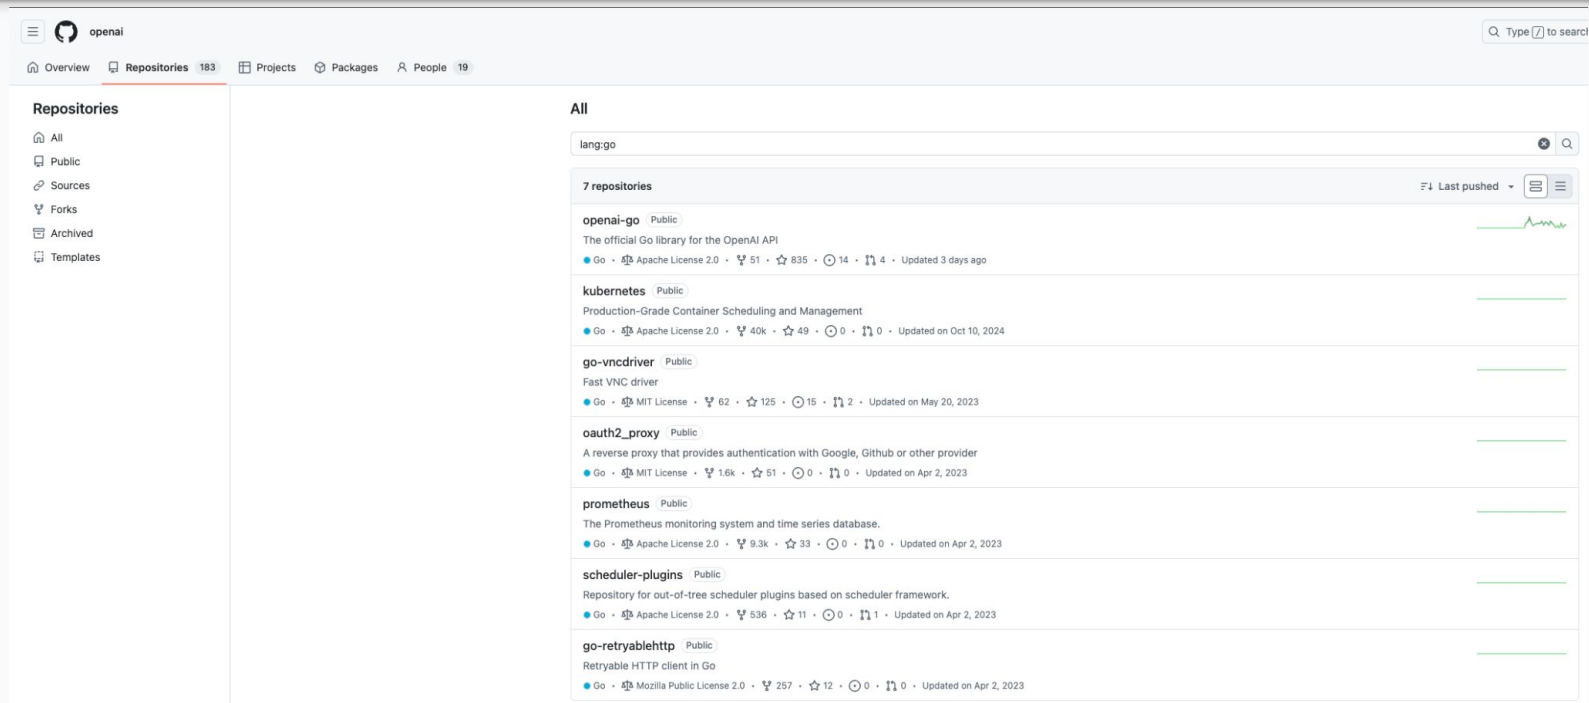
		Representation	
		Explicit	Implicit
Domain Language	Expressive	AaC, Modern automation languages	
	Stringly Typed	Traditional IaC	IfC

Ref: <https://architectelevators.com/cloud/iac-ifc-trends/>

Infrastructure from/with Code vs K8S ecosystem?

- Pulumi - well established project with its own ecosystem
 - Can write custom logic to support non-standard resources? ✓
 - Can write custom workflows? ✓
 - Good base of common cloud resources? ✓
 - Are there standard interfaces? ✗
- Winglang - new language to build infra with code
 - Can write custom logic to support non-standard resources? ✓
 - Can write custom workflows? ✓ (but at what cost?)
 - Good base of common cloud resources? ✗
 - Are there standard interfaces? ✗

OpenAI

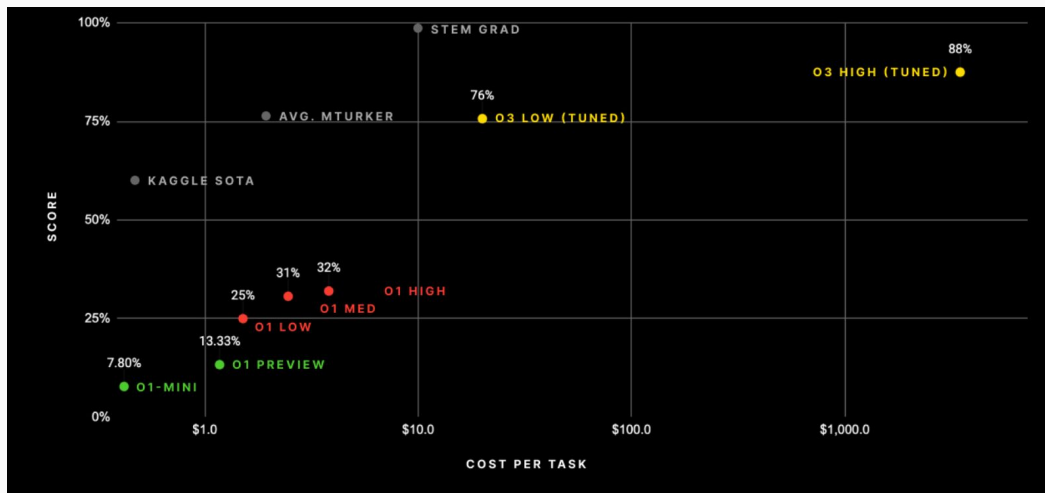
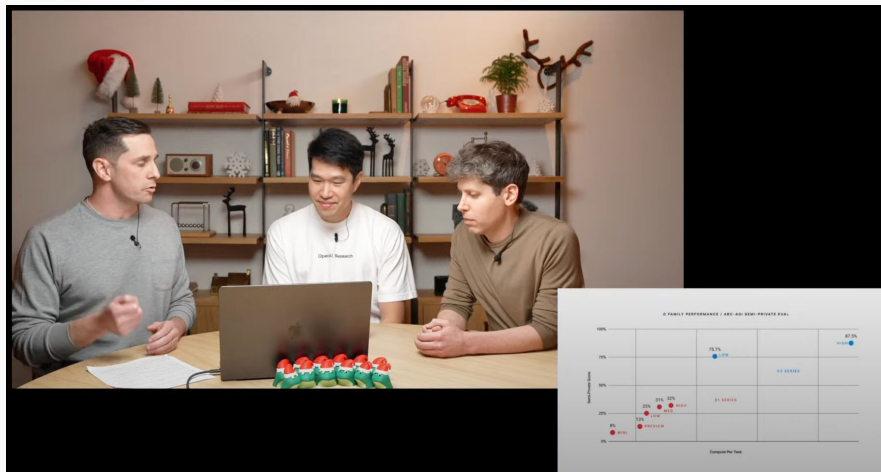


The screenshot displays the OpenAI GitHub profile page. The header includes the OpenAI logo and a search bar. The navigation bar shows 'Overview', 'Repositories 183', 'Projects', 'Packages', and 'People 19'. The left sidebar lists repository filters: 'All', 'Public', 'Sources', 'Forks', 'Archived', and 'Templates'. The main content area, titled 'All', shows a search filter 'lang:go' and a list of 7 repositories. Each repository entry includes its name, description, license, star/fork counts, and the last update date.

Repository	Description	License	Stars	Forks	Updated
openai-go	The official Go library for the OpenAI API	Apache License 2.0	51	835	Updated 3 days ago
kubernetes	Production-Grade Container Scheduling and Management	Apache License 2.0	40k	49	Updated on Oct 10, 2024
go-vncdriver	Fast VNC driver	MIT License	62	125	Updated on May 20, 2023
oauth2_proxy	A reverse proxy that provides authentication with Google, Github or other provider	MIT License	1.6k	51	Updated on Apr 2, 2023
prometheus	The Prometheus monitoring system and time series database.	Apache License 2.0	9.3k	33	Updated on Apr 2, 2023
scheduler-plugins	Repository for out-of-tree scheduler plugins based on scheduler framework.	Apache License 2.0	536	11	Updated on Apr 2, 2023
go-retryablehttp	Retryable HTTP client in Go	Mozilla Public License 2.0	257	12	Updated on Apr 2, 2023

OpenAI

Will AI replace SWEs? There's a lot of marketing and optimism in the field.



Refs:
<https://www.youtube.com/watch?v=SKBG1sodvIU>
<https://techcrunch.com/2024/12/23/openais-o3-suggests-ai-models-are-scaling-in-new-ways-but-so-are-the-costs/>

kthxbye!

Questions?