

# Go 1.24 updates

Patryk Orwat



# Lots of changes in different areas

## Go 1.24 Release Notes

### Table of Contents

Introduction to Go 1.24

Changes to the language

#### Tools

Go command

Cgo

Objdump

Vet

GOCACHEPROG

Runtime

Compiler

Linker

Bootstrap

Standard library

Directory-limited filesystem access

New benchmark function

Improved finalizers

New weak package

New crypto/mlkem package

New crypto/hkdf, crypto/pbkdf2, and crypto/sha3 packages

FIPS 140-3 compliance

New experimental testing/syncctest package

Minor changes to the library

#### Ports

Linux

Darwin

WebAssembly

Windows

More interactive changes: <https://antonz.org/go-1-24/>

# Comment 1: Weak pointers



# Weak pointers - Java good analogy?

From Java tutorial:

*WeakReferences can be used, for example, to store some information related to an object until the object get finalised. To do this you can implement a Map in which the **keys** are wrapped in a WeakReference. As soon as GC will reclaim the key object, you can remove the value as well.*

*Of course it can be done also using some notification mechanism, but using GC will be more robust and efficient. As example you can look at `java.util.WeakHashMap` but it is not thread-safe.*

Source: <https://medium.com/@ramtop/weak-soft-and-phantom-references-in-java-and-why-they-matter-c04bfc9dc792>

# Weak pointers - good for cache? NO!

From Java tutorial:

*Why the WeakHashMap doesn't work for caching? First of all it wouldn't work anyway because it uses soft references for the keys and not for the map values. But additional to that, the garbage collector aggressively reclaims the memory that is referenced only by weak references. It means that once you lose the last strong reference to an object that is working as a key in a WeakHashMap, the garbage collector will soon reclaim that map entry.*

Source: <https://web.archive.org/web/20150403082405/http://www.codeinstructions.com/2008/09/weakhashmap-is-not-cache-understanding.html>

# Weak pointers - 3D rendering story time

Also, good story

<https://stackoverflow.com/a/48048620>

# Weak pointers

And another interesting case - gauges:

*All of the different forms of creating a gauge maintain only a weak reference to the object being observed, so as not to prevent garbage collection of the object.*

Source: <https://docs.micrometer.io/micrometer/reference/concepts/gauges.html>

# Weak pointers - summary

So when to use weak pointers?

In principle, where you want to store additional information to object that you don't own.

In other cases, it depends but usual answer is NO - weak pointers are only for corner cases.



# Weak pointers - good for additional data cache? Maybe..

Yes, you can walk around although cache keys cleanup is just bad...

Maybe [addCleanup](#) would help? If yes, what's the point of weak pointers...

A scheduled job that would iterate over map and delete entries to nil pointers.

That can allow to remove data quickly once strong references are removed.

```
// Cache represents a thread-safe cache with weak pointers.
type Cache[K comparable, V any] struct {
    mu sync.Mutex
    items map[K]weak.Pointer[V] // Weak pointers to cached objects
}

// NewCache creates a new generic Cache instance.
func NewCache[K comparable, V any]() *Cache[K, V] {
    return &Cache[K, V]{
        items: make(map[K]weak.Pointer[V]),
    }
}

// Get retrieves an item from the cache, if it's still alive.
func (c *Cache[K, V]) Get(key K) (*V, bool) {
    c.mu.Lock()
    defer c.mu.Unlock()

    // Retrieve the weak pointer for the given key
    ptr, exists := c.items[key]
    if !exists {
        return nil, false
    }

    // Attempt to dereference the weak pointer
    val := ptr.Value()
    if val == nil {
        // Object has been reclaimed by the garbage collector
        delete(c.items, key)
        return nil, false
    }

    return val, true
}
```

```
// Set adds an item to the cache.
func (c *Cache[K, V]) Set(key K, value V) {
    c.mu.Lock()
    defer c.mu.Unlock()

    // Create a weak pointer to the value
    c.items[key] = weak.Make(&value)
}

func main() {
    // Create a cache with string keys and string values
    cache := NewCache[string, string]()

    // Add an object to the cache
    data := "cached data"
    cache.Set("key1", data)

    // Retrieve it
    if val, ok := cache.Get("key1"); ok {
        fmt.Println("Cache hit:", *val)
    } else {
        fmt.Println("Cache miss")
    }

    // Simulate losing the strong reference
    data = ""
    runtime.GC() // Force garbage collection

    // Try to retrieve it again
    time.Sleep(1 * time.Second)
    if val, ok := cache.Get("key1"); ok {
        fmt.Println("Cache hit:", *val)
    } else {
        fmt.Println("Cache miss")
    }
}
```

# Comment 2: improved finalizer addCleanup

```
type Blob []byte 3 usages new *

func (b Blob) String() string { new *
    return fmt.Sprintf( format: "Blob(%d KB)", len(b)/1024)
}

func newBlob(size int) *Blob { 1 usage new *
    b := make([]byte, size*1024)
    for i := range size {
        b[i] = byte(i) % 255
    }
    return (*Blob>(&b)
}
```

```
func main() { 1 Patryk Orwat *
    b := newBlob( size: 1000)
    now := time.Now()
    // Register a cleanup function to run
    // when the object is no longer reachable.
    runtime.AddCleanup(b, cleanup, now)

    time.Sleep(10 * time.Millisecond)
    b = nil
    runtime.GC()
    time.Sleep(10 * time.Millisecond)
}

func cleanup(created time.Time) { 1 usage new *
    fmt.Printf(
        format: "object is cleaned up! lifetime = %dms\n",
        time.Since(created)/time.Millisecond,
    )
}
```

# Comment 4: FIPS-140-3 compliance

The Go Cryptographic Module is a collection of standard library Go packages under `crypto/internal/fips140/...` that **implement FIPS 140-3 approved algorithms**.

Public API packages such as `crypto/ecdsa` and `crypto/rand` transparently use the Go Cryptographic Module to implement FIPS 140-3 algorithms.

Go Cryptographic Module version v1.0.0 **is currently under test** with a CMVP-accredited laboratory.

Source: <https://go.dev/doc/security/fips140>

# Comment 5: Swiss table

Go 1.23:

Lookup time: *318.447458ms*  
Insertion time: *103.009625ms*  
Deletion time: *36.222416ms*

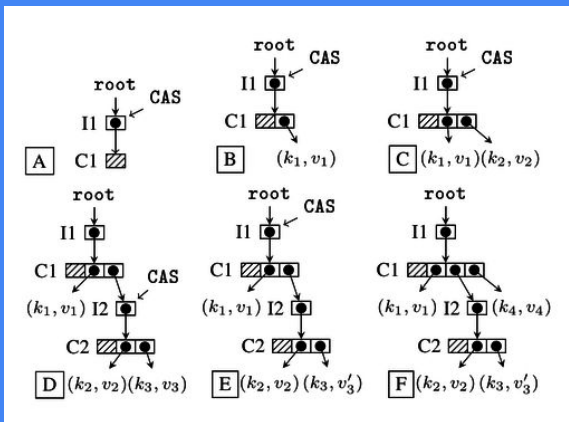
Go 1.24

Lookup time: *237.979625ms*  
Insertion time: *60.243833ms*  
Deletion time: *58.681917ms*

Source: <https://www.bytesizego.com/blog/go-124-swiss-table-maps>

```
func main() {  
    // Create a large map  
    m := make(map[int]int, 1_000_000)  
    for i := 0; i < 1_000_000; i++ {  
        m[i] = i  
    }  
  
    // Measure lookup performance  
    start := time.Now()  
    for i := 0; i < 10_000_000; i++ {  
        _ = m[i%1_000_000]  
    }  
    fmt.Printf("Lookup time: %v\n", time.Since(start))  
  
    // Measure insertion performance  
    start = time.Now()  
    for i := 1_000_000; i < 2_000_000; i++ {  
        m[i] = i  
    }  
    fmt.Printf("Insertion time: %v\n", time.Since(start))  
  
    // Measure deletion performance  
    start = time.Now()  
    for i := 0; i < 1_000_000; i++ {  
        delete(m, i)  
    }  
    fmt.Printf("Deletion time: %v\n", time.Since(start))  
}
```

# Comment 6: sync.map is Ctrie



Source: <https://en.wikipedia.org/wiki/Ctrie>

## Before Go 1.24

### Underlying Structure

```
clike Copy

type Map struct {
    mu    Mutex
    read  atomic.Value // readOnly
    dirty map[interface{}]*entry
    misses int
}

type readOnly struct {
    m    map[interface{}]*entry
    amended bool
}
```

- **mu** : A mutex used to protect access to **read** and **dirty** .
- **read** : A read-only data structure supporting concurrent reads using atomic operations. It stores a **readOnly** structure, which is a native map. The **amended** attribute marks whether the **read** and **dirty** data are consistent.
- **dirty** : A native map for reading and writing data, requiring locking to ensure data security.
- **misses** : A counter tracking how many times the read operation fails.

### Entry Structure

```
clike Copy

type entry struct {
    p unsafe.Pointer // *interface{}
}
```

- It contains a pointer **p** that points to the value stored for the element (key).

Source: <https://reliasoftware.com/blog/go-sync-map>

Figure 1 consists of six sub-diagrams labeled A through F, illustrating the evolution of a binary tree structure. Each diagram shows a hierarchy of nodes: root, II, C1, I2, and C2. Nodes are represented by boxes: white for root, II, and I2; hatched for C1 and C2. Black dots represent nodes in the CAS (Current Active Set). Arrows indicate the flow of information or updates.

- Diagram A:** Shows the root node pointing to II and C1. II points to C1. C1 contains a black dot. The label "CAS" points to the black dot in C1.
- Diagram B:** Similar to A, but C1 contains two black dots. The label "(k<sub>1</sub>, v<sub>1</sub>)" is below C1.
- Diagram C:** Similar to B, but I2 is added as a child of C1. I2 contains a black dot. The label "(k<sub>1</sub>, v<sub>1</sub>)(k<sub>2</sub>, v<sub>2</sub>)" is below I2.
- Diagram D:** Similar to C, but C2 is added as a child of I2. C2 contains two black dots. The label "(k<sub>2</sub>, v<sub>2</sub>)(k<sub>3</sub>, v<sub>3</sub>)" is below C2.
- Diagram E:** Similar to D, but C1 contains three black dots. The label "(k<sub>2</sub>, v<sub>2</sub>)(k<sub>3</sub>, v<sub>3</sub>)" is below C1.
- Diagram F:** Similar to E, but I2 contains two black dots. The label "(k<sub>2</sub>, v<sub>2</sub>)(k<sub>3</sub>, v<sub>3</sub>)" is below I2.

Source: <https://en.wikipedia.org/wiki/Ctrie>

	before		after	
	sec/op		sec/op	
				vs base
MapLoadMostlyHits	7.870n ±	1%	8.415n ±	3% +6.93%
MapLoadMostlyMisses	7.210n ±	1%	5.314n ±	2% -26.28%
MapLoadOrStoreBalanced	360.10n ±	18%	71.78n ±	2% -80.07%
MapLoadOrStoreUnique	707.2n ±	18%	135.2n ±	4% -80.88%
MapLoadOrStoreCollision	5.089n ±	201%	3.963n ±	1% -22.11%
MapLoadAndDeleteBalanced	17.045n ±	64%	5.280n ±	1% -69.02%
MapLoadAndDeleteUnique	14.250n ±	57%	6.452n ±	1% ~
MapLoadAndDeleteCollision	19.34n ±	39%	23.31n ±	27% ~
MapRange	3.055μ ±	3%	1.918μ ±	2% -37.23%
MapAdversarialAlloc	245.30n ±	6%	14.90n ±	23% -93.92%
MapAdversarialDelete	143.550n ±	2%	8.184n ±	1% -94.30%
MapDeleteCollision	9.199n ±	65%	3.165n ±	1% -65.59%
MapSwapCollision	164.7n ±	7%	108.7n ±	36% -34.01%
MapSwapMostlyHits	33.12n ±	15%	35.79n ±	9% ~
MapSwapMostlyMisses	604.5n ±	5%	280.2n ±	7% -53.64%
MapCompareAndSwapCollision	96.02n ±	40%	69.93n ±	24% -27.17%
MapCompareAndSwapNoExistingKey	6.345n ±	1%	6.202n ±	1% -2.24%
MapCompareAndSwapValueNotEqual	6.121n ±	3%	5.564n ±	4% -9.09%
MapCompareAndSwapMostlyHits	44.21n ±	13%	43.46n ±	11% ~
MapCompareAndSwapMostlyMisses	33.51n ±	6%	13.51n ±	5% -59.70%
MapCompareAndDeleteCollision	27.85n ±	104%	31.02n ±	26% ~
MapCompareAndDeleteMostlyHits	50.43n ±	33%	109.45n ±	8% +117.03%
MapCompareAndDeleteMostlyMisses	27.17n ±	7%	11.37n ±	3% -58.14%
MapClear	300.2n ±	5%	124.2n ±	8% -58.64%
geomean	50.38n		25.79n	

The load-hit case (`MapLoadMostlyHits`) is slightly slower due to Swiss Tables improving the performance of the old `sync.Map`. Some benchmarks show a seemingly large slowdown, but that's mainly due to the fact that the new implementation shrinks promptly (as elements are deleted from the map), whereas the old one shrinks in generations (the dirty map needs to be promoted).

Source: <https://antonz.org/go-1-24/#concurrent-hash-trie-map>

# Demo

Feat 1: Tool dependencies

Feat 2: main module version

Feat 3&4: working dir in test & test context